

Overview of Computer Architecture

Processors and Execution

Carl Henrik Ek - carlhenrik.ek@bristol.ac.uk

November 8, 2019

<http://carlhenrik.com>

Definition (Computer Science)

The systematic study (science) of algorithms

Definition (Algorithms)

an **ordered** set of **unambiguous**, **executable** steps that defines a **terminating** process

Computing



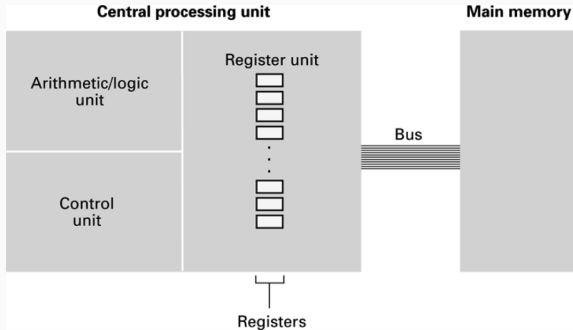
Computer Architecture



Execution

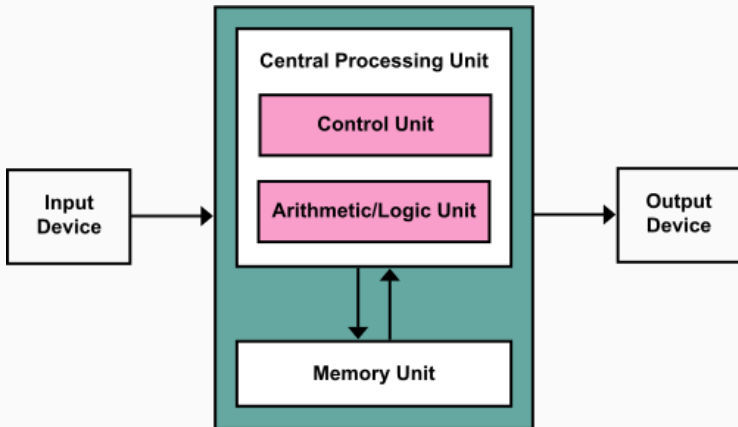


Computer Architecture

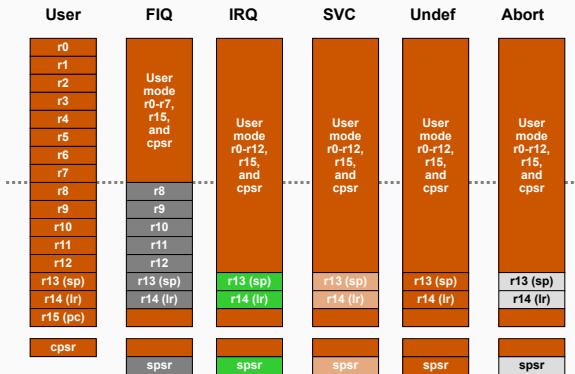


RAM as an array of bytes

Content:	FF	00	57	92	B3	8A	10	46	DC
Address:	000 000 000	000 000 001	000 000 002	000 000 003	000 000 004	000 000 005	...	134 217 725	134 217 726	134 217 727



Register



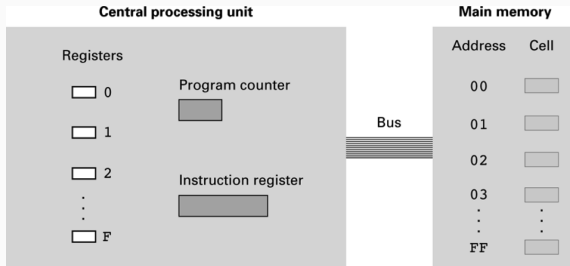
Program Counter

- Address to memory to fetch instruction to execute
- Start position


Instruction Register

- Stores instruction


Execution



Execution

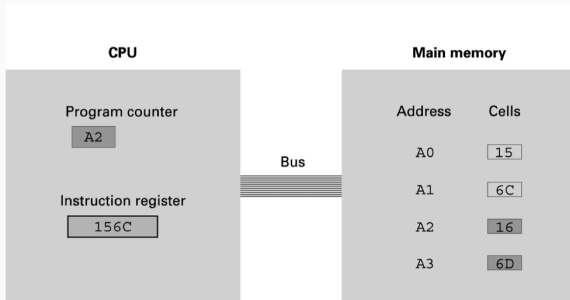


Execution



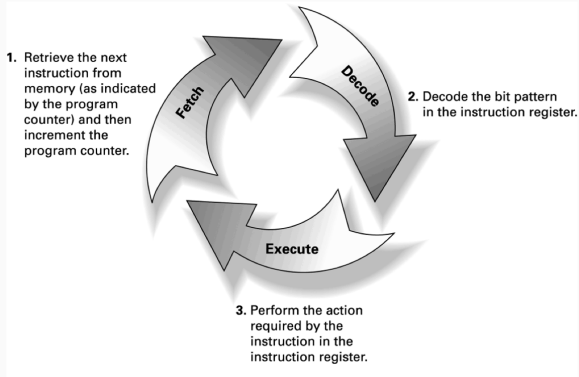
- a. At the beginning of the fetch step the instruction starting at address A0 is retrieved from memory and placed in the instruction register.

Execution



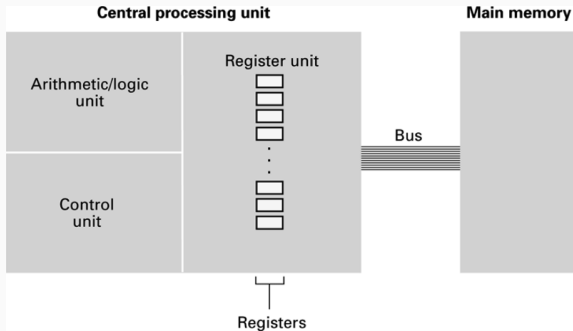
b. Then the program counter is incremented so that it points to the next instruction.

Execution




Processors


Computer Architecture



State Machine



Status Register



Copies of the ALU status flags (latched if the instruction has the "S" bit set).

* Condition Code Flags

N = Negative result from ALU flag.

Z = Zero result from ALU flag.

C = ALU operation **C**arried out

V = ALU operation overflowed

* **Mode Bits**

M[4:0] define the processor mode.

* Interrupt Disable bits.

$I = 1$, disables the IRQ.

F = 1, disables the FIQ.

* T Bit (Architecture v4T only)

T = 0, Processor in ARM state

T = 1, Processor in Thumb state

Flags

	Logical Instruction	Arithmetic Instruction
<u>Flag</u>		
Negative (N='1')	No meaning	Bit 31 of the result has been set Indicates a negative number in signed operations
Zero (Z='1')	Result is all zeroes	Result of operation was zero
Carry (C='1')	After Shift operation '1' was left in carry flag	Result was greater than 32 bits
oVerflow (V='1')	No meaning	Result was greater than 31 bits Indicates a possible corruption of the sign bit in signed numbers

Conditions

3 1	3 0	2 9	2 8	2 7	2 6	2 5	2 4	2 3	2 2	2 1	2 0	1 9	1 8	1 7	1 6	1 5	1 4	1 3	1 2	1 1	1 0	9	8	7	6	5	4	3	2	1	0	Instruction Type
Condition			0	0	1	OPCODE						S	Rn				Rs				OPERAND-2											Data processing

0000 = EQ - Z set (equal)

0001 = NE - Z clear (not equal)

0010 = HS / CS - C set (unsigned higher or same)

0011 = LO / CC - C clear (unsigned lower)

0100 = MI -N set (negative)

0101 = PL - N clear (positive or zero)

0110 = VS - V set (overflow)

0111 = VC - V clear (no overflow)

1000 = HI - C set and Z clear (unsigned higher)

1001 = LS - C clear or Z (set unsigned lower or same)

1010 = GE - N set and V set, or N clear and V clear (>or =)

1011 = LT - N set and V clear, or N clear and V set (>)

1100 = GT - Z clear, and either N set and V set, or N clear and V set (>)

1101 = LE - Z set, or N set and V clear, or N clear and V set (<, or =)

1110 = AL - always

1111 = NV - reserved.

- OP-Code
 - Decides operation
- Operand
 - OP-Code dependent
 - "Parameters"
- Each instruction on ARM 32 bits

Instruction ARM

3 1	3 0	2 9	2 8	2 7	2 6	2 5	2 4	2 3	2 2	2 1	2 0	1 9	1 8	1 7	1 6	1 5	1 4	1 3	1 2	1 1	1 0	9	8	7	6	5	4	3	2	1	0	Instruction Type											
Condition				0	0	I	OPCODE				S	Rn				Rs				OPERAND-2								Data processing															
Condition				0	0	0	0	0	0	A	S	Rd				Rn				Rs				1	0	0	1	Rm				Multiply											
Condition				0	0	0	0	0	1	U	A	S	Rd HIGH				Rd LOW				Rs				1	0	0	1	Rm				Long Multiply										
Condition				0	0	0	1	0	B	0	0	Rn				Rd				0	0	0	0	1	0	0	1	Rm				Swap											
Condition				0	1	I	P	U	B	W	L	Rn				Rd				OFFSET								Load/Store - Byte/Word															
Condition				1	0	0	P	U	B	W	L	Rn				REGISTER LIST								Load/Store Multiple																			
Condition				0	0	0	P	U	1	W	L	Rn				Rd				OFFSET 1				1	S	H	1	OFFSET 2				Halfword Transfer Imm Off											
Condition				0	0	0	P	U	0	W	L	Rn				Rd				0	0	0	0	1	S	H	1	Rm				Halfword Transfer Reg Off											
Condition				1	0	1	L	BRANCH OFFSET																												Branch							
Condition				0	0	0	1	0	0	1	0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	0	0	0	1	Rn				Branch Exchange									
Condition				1	1	0	P	U	N	W	L	Rn				CRd				CPNum				OFFSET								COPROCESSOR DATA XFER											
Condition				1	1	1	0	Op-1				CRn				CRd				CPNum				OP-2				0	CRm				COPROCESSOR DATA OP										
Condition								OP-1				L	CRn				Rd				CPNum				OP-2				1	CRm				COPROCESSOR REG XFER									
Condition				1	1	1	1	SWI NUMBER																												Software Interrupt							



- Data Transfer
 - LDR, STR
- Flow Control
 - B, CMP
- Arithmetic
 - ADD, MUL

4.9.8 Assembler syntax

<LDR|STR>{cond}{B}{T} Rd,<Address>

where:

- LDR load from memory into a register
- STR store from a register into memory
- {cond} two-character condition mnemonic. See **Table 4-2: Condition code summary** on page 4-5.
- {B} If B is present then byte transfer, otherwise word transfer
- {T} If T is present the W bit will be set in a post-indexed instruction, forcing non-privileged mode for the transfer cycle. T is not allowed when a pre-indexed addressing mode is specified or implied.
- Rd is an expression evaluating to a valid register number.
- Rn and Rm are expressions evaluating to a register number. If Rn is R15 then the assembler will subtract 8 from the offset value to allow for pipelining. In this case base write-back should not be specified.

<Address> can be:

- 1 An expression which generates an address:

<expression>

The assembler will attempt to generate an instruction using the PC as a base and a corrected immediate offset to address the location given by evaluating the expression. This will be a PC relative, pre-indexed address. If the address is out of range, an error will be generated.

- 2 A pre-indexed addressing specification:

[Rn] offset of zero

[Rn,<#expression>]{!} offset of <expression> bytes

[Rn,{+/-}Rm{,<shift>}]!} offset of +/- contents of index register, shifted by <shift>

- 3 A post-indexed addressing specification:

[Rn],<#expression> offset of <expression> bytes

[Rn,{+/-}Rm{,<shift>}] offset of +/- contents of index register, shifted as by <shift>.

<shift> general shift operation (see data processing instructions) but you cannot specify the shift amount by a register.

`ldr r0,[r1,#4]` Load word addressed by $R1+4$.

`str r0,[r1],#4` Store R0 to word addressed by R1. Increment R1 by 4.

`ldr r0,[r1,#4]!` Load word addressed by $R1+4$. Increment R1 by 4.

`ldr r0,=label` Load address of label `label` into R0

Code

```
_start:
```

```
    add    r0, r1
```

```
    add    r0, #4
```

```
    add    r0, [r1, #3]
```

4.7 Multiply and Multiply-Accumulate (MUL, MLA)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-12: Multiply instructions](#).

The multiply and multiply-accumulate instructions use an 8 bit Booth's algorithm to perform integer multiplication.




Figure 4-12: Multiply instructions

The multiply form of the instruction gives $Rd = Rm * Rs$. Rn is ignored, and should be set to zero for compatibility with possible future upgrades to the instruction set.

The multiply-accumulate form gives $Rd = Rm * Rs + Rn$, which can save an explicit ADD instruction in some circumstances.

Both forms of the instruction work on operands which may be considered as signed (2's complement) or unsigned integers.

The results of a signed multiply and of an unsigned multiply of 32 bit operands differ only in the upper 32 bits - the low 32 bits of the signed and unsigned results are identical. As these instructions only produce the low 32 bits of a multiply, they can be used for both signed and unsigned multiplies.

For example consider the multiplication of the operands:

Operand A Operand B Result
0xFFFFFFFF6 0x0000010xFFFFFFFF38

If the operands are interpreted as signed

Operand A has the value -10, operand B has the value 20, and the result is -200 which is correctly represented as 0xFFFFFFFF38

If the operands are interpreted as unsigned

Operand A has the value 4294967286, operand B has the value 20 and the result is 85899345720, which is represented as 0x13FFFFFF38, so the least significant 32 bits are 0xFFFFFFFF38.

4.7.1 Operand restrictions

The destination register Rd must not be the same as the operand register Rm. R15 must not be used as an operand or as the destination register.

4.4 Branch and Branch with Link (B, BL)

The instruction is only executed if the condition is true. The various conditions are defined **Table 4-2: Condition code summary** on page 4-5. The instruction encoding is shown in **Figure 4-3: Branch instructions**, below.




Figure 4-3: Branch instructions

Branch instructions contain a signed 2's complement 24 bit offset. This is shifted left two bits, sign extended to 32 bits, and added to the PC. The instruction can therefore specify a branch of +/- 32Mbytes. The branch offset must take account of the prefetch operation, which causes the PC to be 2 words (8 bytes) ahead of the current instruction. Branches beyond +/- 32Mbytes must use an offset or absolute destination which has been previously loaded into a register. In this case the PC should be manually saved in R14 if a Branch with Link type operation is required.

4.4.1 The link bit

Branch with Link (BL) writes the old PC into the link register (R14) of the current bank. The PC value written into R14 is adjusted to allow for the prefetch, and contains the address of the instruction following the branch and link instruction. Note that the CPSR is not saved with the PC and R14[1:0] are always cleared.

To return from a routine called by Branch with Link use MOV PC,R14 if the link register is still valid or LDM Rn!,{..PC} if the link register has been saved onto a stack pointed to by Rn.

4.4.2 Instruction cycle times

Branch and Branch with Link instructions take 2S + 1N incremental cycles, where S and N are as defined in **6.2 Cycle Types** on page 6-3.

Code

```
.section          .text
.align            2
.global           _start

_start:

    ldr           r0, =matrix

    ldr           r1, [r0]           @ a
    ldr           r2, [r0, #12]     @ d
    mul           r7, r2, r1        @ a*d

    ldr           r1, [r0, #4]      @ b
    ldr           r2, [r0, #8]      @ c
    mul           r6, r2, r1        @ b*c
```


Summary

Summary

- This was just a very quick intro to CPUs
- You learn this by doing
- Lab on Monday we will play with these things