



Overview of Computer Architecture

Processors and Execution

Carl Henrik Ek - carlhenrik.ek@bristol.ac.uk

November 8, 2019

<http://carlhenrik.com>

Computer Science

Definition (Computer Science)

The systematic study (science) of algorithms

Definition (Algorithms)

an **ordered** set of **unambiguous**, **executable** steps that defines a **terminating** process

Computing



Computer Architecture



Sinclair Spectrum



Commodore 64





Archimedes



ARM

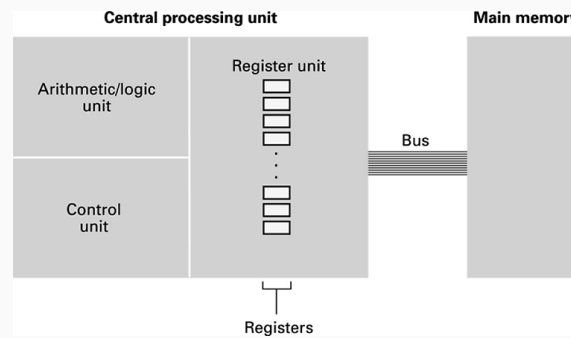
arm

Execution

Start



Computer Architecture




Memory

RAM as an array of bytes

Content:	FF	00	57	92	B3	8A	10	46	DC
Address:	000 000 000	000 000 001	000 000 002	000 000 003	000 000 004	000 000 005	134 217 725	134 217 726	134 217 727

CPU



Register

r0					
r1					
r2					
r3					
r4					
r5					
r6					
r7					
r8					
r9					
r10					
r11					
r12					
r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)	r13 (sp)
r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)	r14 (lr)
r15 (pc)					
cpsr					
	spsr	spsr	spsr	spsr	spsr

Program Counter

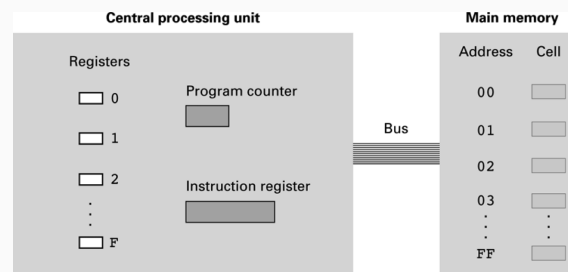
Program Counter

- Address to memory to fetch instruction to execute
- Start position

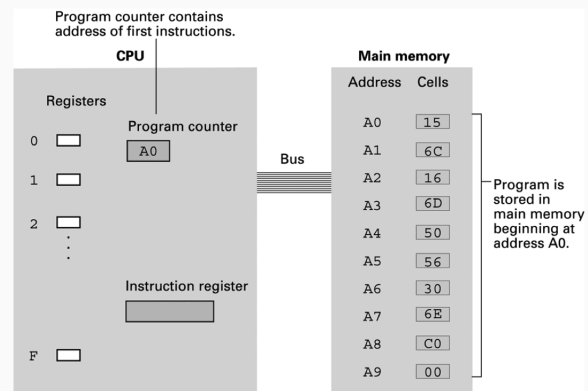
Instruction Register

- Stores instruction


Execution



Execution

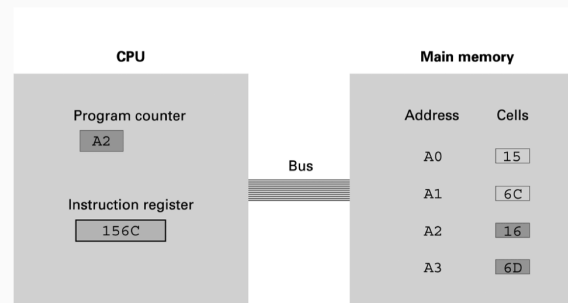


Execution



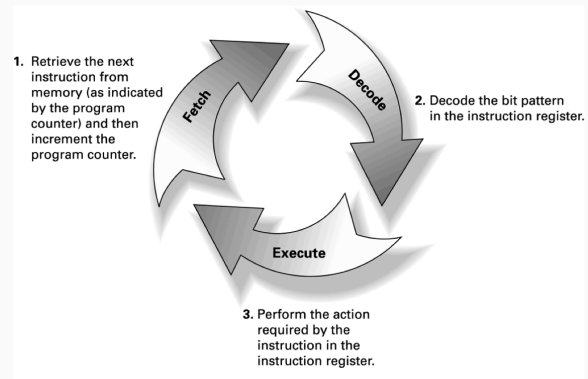
- a. At the beginning of the fetch step the instruction starting at address A0 is retrieved from memory and placed in the instruction register.

Execution



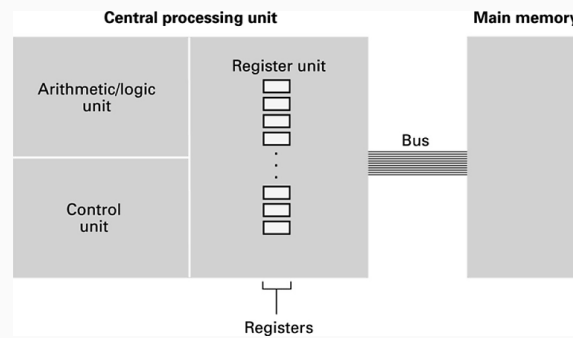
b. Then the program counter is incremented so that it points to the next instruction.

Execution




Processors

Computer Architecture



State Machine



Status Register



Copies of the ALU status flags (latched if the instruction has the "S" bit set).

* **Condition Code Flags**

N = Negative result from ALU flag.
 Z = Zero result from ALU flag.
 C = ALU operation Carried out
 V = ALU operation oVerflowed

* **Mode Bits**

M[4:0] define the processor mode.

* **Interrupt Disable bits.**

I = 1, disables the IRQ.
 F = 1, disables the FIQ.

* **T Bit (Architecture v4T only)**

T = 0, Processor in ARM state
 T = 1, Processor in Thumb state

Flags

	Logical Instruction	Arithmetic Instruction
<u>Flag</u>		
Negative (N='1')	No meaning	Bit 31 of the result has been set Indicates a negative number in signed operations
Zero (Z='1')	Result is all zeroes	Result of operation was zero
Carry (C='1')	After Shift operation '1' was left in carry flag	Result was greater than 32 bits
oVerflow (V='1')	No meaning	Result was greater than 31 bits Indicates a possible corruption of the sign bit in signed numbers

Conditions

3 1	2 0	2 9	2 8	2 7	2 6	2 5	2 4	2 3	2 2	2 1	2 0	1 1	1 0	1 1	1 0	1 1	1 0	1 1	1 0	1 1	1 0	0 0	0 1	0 2	0 3	0 4	0 5	0 6	0 7	0 8	0 9	0 10	0 11	0 12	0 13	0 14	0 15	Instruction Type
Condition	0	0	1																																			Data processing

0000 = EQ - Z set (equal)	1001 = LS - C clear or Z (set unsigned lower or same)
0001 = NE - Z clear (not equal)	1010 = GE - N set and V set, or N clear and V clear (>or =)
0010 = HS / CS - C set (unsigned higher or same)	1011 = LT - N set and V clear, or N clear and V set (>)
0011 = LO / CC - C clear (unsigned lower)	1100 = GT - Z clear, and either N set and V set, or N clear and V set (>)
0100 = MI - N set (negative)	1101 = LE - Z set, or N set and V clear, or N clear and V set (<, or =)
0101 = PL - N clear (positive or zero)	1110 = AL - always
0110 = VS - V set (overflow)	1111 = NV - reserved.
0111 = VC - V clear (no overflow)	
1000 = HI - C set and Z clear (unsigned higher)	


Instructions

- OP-Code
 - Decides operation
- Operand
 - OP-Code dependent
 - "Parameters"
- Each instruction on ARM 32 bits

Instruction ARM

31	30	29	28	27	26	25	24	23	22	21	20	19	18	17	16	15	14	13	12	11	10	9	8	7	6	5	4	3	2	1	0	Instruction Type		
Condition	0	0	1	OPCODE					S	Rn			Rs			OPERAND-2					Data processing													
Condition	0	0	0	0	0	0	A	S	Rd			Rn			Rs			1	0	0	1	Rm							Multiply					
Condition	0	0	0	0	1	U	A	S	Rd HIGH			Rd LOW			Rs			1	0	0	1	Rm							Long Multiply					
Condition	0	0	0	1	0	B	0	0	Rn			Rd			0	0	0	0	1	0	0	1	Rm							Swap				
Condition	0	1	I	P	U	B	W	L	Rn			Rd			OFFSET					Load/Store - Byte/Word														
Condition	1	0	0	P	U	B	W	L	Rn			REGISTER LIST					Load/Store Multiple																	
Condition	0	0	0	P	U	1	W	L	Rn			Rd			OFFSET 1			1	S	H	1	OFFSET 2			Halfword Transfer Imm Off									
Condition	0	0	0	P	U	0	W	L	Rn			Rd			0	0	0	0	1	S	H	1	Rm			Halfword Transfer Reg Off								
Condition	1	0	1	L	BRANCH OFFSET														Branch															
Condition	0	0	0	1	0	0	1	0	1	1	1	1	1	1	1	1	1	1	1	0	0	0	1	Rn							Branch Exchange			
Condition	1	1	0	P	U	N	W	L	Rn			CRd			CPN um			OFFSET					COPROCESSOR DATA XFER											
Condition	1	1	1	0	Op-1			CRn			CRd			CPN um			OP-2			0	CRm			COPROCESSOR DATA OP										
Condition				OP-1			L	CRn			Rd			CPN um			OP-2			1	CRm			COPROCESSOR REG XFER										
Condition	1	1	1	1	SWI NUMBER																Software Interrupt													

Reference Manual



Instructions

- Data Transfer
 - LDR, STR
- Flow Control
 - B, CMP
- Arithmetic
 - ADD, MUL

LDR/STR

4.9 Single Data Transfer (LDR, STR)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-14: Single data transfer instructions](#).

The single data transfer instructions are used to load or store single bytes or words of data. The memory address used in the transfer is calculated by adding an offset to or subtracting an offset from a base register.

The result of this calculation may be written back into the base register if auto-indexing is required.

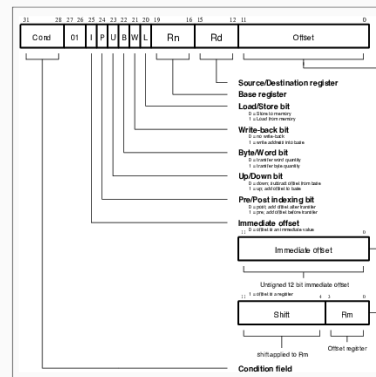


Figure 4-14: Single data transfer instructions

Addressing Modes

`ldr r0,[r1,#4]` Load word addressed by $R1+4$.

`str r0,[r1],#4` Store $R0$ to word addressed by $R1$. Increment $R1$ by 4.

`ldr r0,[r1,#4]!` Load word addressed by $R1+4$. Increment $R1$ by 4.

`ldr r0,=label` Load address of label `label` into $R0$

ADD

Code

```
_start:  
    add    r0, r1  
    add    r0, #4  
    add    r0, [r1, #3]
```

MUL

4.7 Multiply and Multiply-Accumulate (MUL, MLA)

The instruction is only executed if the condition is true. The various conditions are defined in [Table 4-2: Condition code summary](#) on page 4-5. The instruction encoding is shown in [Figure 4-12: Multiply instructions](#).

The multiply and multiply-accumulate instructions use an 8-bit Booth's algorithm to perform integer multiplication.




Figure 4-12: Multiply instructions

The multiply form of the instruction gives $Rd = Rn \times Rm$. Rn is ignored, and should be set to zero for compatibility with possible future upgrades to the instruction set. The multiply-accumulate form gives $Rd = Rn + Rm$, which can save an explicit ADD instruction in some circumstances.

Both forms of the instruction work on operands which may be considered as signed (2's complement) or unsigned integers.

The results of a signed multiply and of an unsigned multiply of 32-bit operands differ only in the upper 32 bits – the low 32 bits of the signed and unsigned results are identical. As these instructions only produce the low 32 bits of a multiply, they can be used for both signed and unsigned multiplies.

For example consider the multiplication of the operands:

Operand A: operand B: result
0xFFFFFFFF 0x00000010 0xFFFFFFFF38

If the operands are interpreted as signed

Operand A has the value -10, operand B has the value 20, and the result is -200 which is correctly represented as 0xFFFFFFFF38.

If the operands are interpreted as unsigned

Operand A has the value 4294967286, operand B has the value 20 and the result is 85899345720, which is represented as 0xFFFFFFFF38, so the least significant 32 bits are 0xFFFFFFFF38.

4.7.1 Operand restrictions

The destination register Rd must not be the same as the operand register Rm . $R15$ must not be used as an operand or as the destination register.

Branch

4.4 Branch and Branch with Link (B, BL)

The instruction is only executed if the condition is true. The various conditions are defined *Table 4-2: Condition code summary* on page 4-5. The instruction encoding is shown in *Figure 4-3: Branch instructions*, below.



Branch instructions contain a signed 2's complement 24 bit offset. This is shifted left two bits, sign extended to 32 bits, and added to the PC. The instruction can therefore specify a branch of +/- 32Mbytes. The branch offset must take account of the prefetch operation, which causes the PC to be 2 words (8 bytes) ahead of the current instruction. Branches beyond +/- 32Mbytes must use an offset or absolute destination which has been previously loaded into a register. In this case the PC should be manually saved in R14 if a Branch with Link type operation is required.

4.4.1 The link bit

Branch with Link (BL) writes the old PC into the link register (R14) of the current bank. The PC value written into R14 is adjusted to allow for the prefetch, and contains the address of the instruction following the branch and link instruction. Note that the CPSR is not saved with the PC and R14[1:0] are always cleared.

To return from a routine called by Branch with Link use MOV PC,R14 if the link register is still valid or LDM Rn!,{PC} if the link register has been saved onto a stack pointed to by Rn.

4.4.2 Instruction cycle times

Branch and Branch with Link instructions take $2S + 1N$ incremental cycles, where S and N are as defined in *6.2 Cycle Types* on page 6-3.

Code

Code

```
.section      .text
.align       2
.global      _start
_start:
    ldr      r0, [matrix]

    ldr      r1, [r0]      @ a
    ldr      r2, [r0, #12] @ d
    mul      r7, r2, r1    @ a*d

    ldr      r1, [r0, #4]  @ b
    ldr      r2, [r0, #8]  @ c
    mul      r6, r2, r1    @ b*c
```

Summary

Summary

- This was just a very quick intro to CPUs
- You learn this by doing
- Lab on Monday we will play with these things